

On the Token Reward Quality of On-Policy Distillation

Anonymous Authors

Abstract

On-policy knowledge distillation trains the student on its own rollouts while the teacher provides token-level soft targets. A common assumption is that the teacher’s supervision is uniformly useful across all token positions. We show this assumption is wrong: teacher signal quality varies dramatically by position. At early positions, the teacher provides an independent assessment of reasoning strategy (entropy 1.07, 42% token-level disagreement with the student); at later positions, the teacher is conditioned on the student’s prefix and largely rubber-stamps the student’s choices (entropy 0.14, 13% disagreement). Training on only the first 100 tokens improves performance by +14.9pp over baseline, while training on the middle or last 100 tokens *decreases* performance below baseline. This reveals that later-position supervision is not merely uninformative—it is actively harmful. We propose *positional distillation*: restricting KD loss to the first N response tokens, yielding a 20–35 $\times$ wall-clock speedup. Experiments across math reasoning (MATH-500), code generation (HumanEval/MBPP), and function calling (BFCL), with multiple teacher scales (1.7B–8B) and model families (Qwen, Gemma), show that positional distillation matches or exceeds full-sequence distillation while eliminating training instability.

1 Introduction

On-policy knowledge distillation (KD) trains the student on its own rollouts while the teacher supplies token-level soft targets [Agarwal et al., 2024, Gu et al., 2024]. From the student’s perspective, this is equivalent to receiving a *dense, token-level reward*: at every position the teacher acts as a reward model, scoring each student-generated token via its conditional distribution [Gu et al., 2024, Lu et al., 2025, Yang et al., 2026]. A natural assumption is that this token reward is uniformly useful across positions—that every token in the student’s trajectory receives an equally valuable signal from the teacher. This paper challenges that assumption.

This token reward is typically understood as the teacher assessing the quality of the student’s answer given the question. But since the reward is computed autoregressively— $p_{\text{teacher}}(y_t \mid x, y_{<t})$, where $y_{<t}$ is the *student’s own prefix*—the interaction between teacher and student naturally decomposes into two regimes. At **early positions**, the prefix is largely determined by the prompt x , so the teacher’s assessment is primarily based on the question itself—an *independent* evaluation. At **later positions**, the teacher is conditioned on hundreds of student-generated tokens; it is no longer evaluating the student based on the question alone, but rather scoring how to continue *this particular* student trajectory—a *dependent* evaluation shaped by the student’s prior choices. These two regimes are fundamentally different in nature. This raises a critical question: *is the teacher’s token reward at later positions still high quality?* We investigate this question and find that the answer is no.

Our analysis reveals several converging lines of evidence. **(1) Signal strength decays with position.** Teacher entropy drops from 1.07 at early positions to 0.14 at late positions; teacher-student agreement rises from 58% to 87%. At later positions, the teacher is effectively rubber-stamping the student’s choices. **(2) Later-position supervision is actively harmful.** Training on the first 100 tokens improves MATH-500 accuracy by +14.9pp over baseline, while training on the middle or last 100 tokens *decreases* accuracy by 1–4pp below baseline. **(3) Full-sequence training degrades over time.** On coding tasks, full-sequence distillation degrades from 40.2% to 26.8% HumanEval—consistent with accumulated exposure to low-quality late-position signals.

These findings motivate **Positional Distillation**: restricting KD loss to the first N response tokens and truncating rollouts accordingly. This is not merely an efficiency trick—it




Figure 1: Teacher signal quality degrades with position (100 MATH-500 problems, 1 trajectory each). (a) Reverse KL divergence $KL(p_s || p_t)$ decays from 2.09 in the first 50 tokens to ~ 0.24 by position 1000. Dashed line shows the token-level log-prob gap, which nearly coincides with the reverse KL. (b) Teacher entropy drops from 0.26 to 0.12 while teacher-student agreement rises from 79% to 93%. At later positions, the teacher is effectively rubber-stamping the student’s choices.

is a principled response to the non-uniform signal quality we identify. Experiments across math reasoning (MATH-500), code generation (HumanEval/MBPP), and function calling (BFCL), with teacher scales from 1.7B to 8B and cross-family validation (Qwen, Gemma), show that positional distillation matches or exceeds full-sequence distillation on all tasks while providing 20–35 $\times$ wall-clock speedup and eliminating training instability.

Our contributions are:

1. **Analysis of teacher signal quality.** We show that in on-policy KD, the teacher’s supervision quality degrades along the trajectory: early tokens carry strong, independent signal about reasoning strategy, while later tokens contribute weak, dependent signal that is actively harmful when trained on (Section 2).
2. **Positional Distillation.** A method that restricts KD loss to the first N response tokens, simultaneously improving quality, stability, and efficiency (Section 3).
3. **Broad evaluation.** Consistent improvements across three tasks, three teacher scales, and two model families, with detailed ablations on position selection and training dynamics (Section 4).

2 Token Reward Quality Analysis by Position

We analyze the quality of the teacher’s per-token supervision signal across positions in student-generated trajectories. All analysis in this section uses Qwen2.5-Math-1.5B as the student and Qwen3-1.7B as the teacher on MATH-500, with 59,936 on-policy trajectories (mean length 386 tokens).

2.1 The Conditional Scoring Problem

In on-policy distillation, the teacher scores position t via $p_{\text{teacher}}(\cdot | x, y_{1:t-1}^{\text{student}})$. This score is *not* the teacher’s unconditional opinion about what token should come next—it is the teacher’s opinion *conditioned on the student’s prefix*. As t increases, the prefix $y_{1:t-1}$ becomes increasingly student-determined, and the teacher’s scoring increasingly reflects “how to continue this student trajectory” rather than “what is the right approach.”

This creates a *phase shift* in the teacher’s role along the trajectory. At early positions, the teacher acts as a **strategic guide**—providing independent judgment on how to approach the problem. At later positions, the teacher degenerates into a **syntax checker**—merely confirming that the student’s execution is locally coherent within the student’s chosen

Problem: In an isosceles right triangle, the altitude to the hypotenuse has length $4\sqrt{2}$. What is the area?

First 100 tokens (strategy): “Let’s start by understanding the problem. In an isosceles right triangle, the two legs are equal, and the hypotenuse is $\sqrt{2}$ times the leg. If we denote each leg by a , then hypotenuse $= a\sqrt{2}$. The altitude to the hypotenuse is half the leg, because the altitude bisects the...”

Last 100 tokens (execution): “...altitude from the right-angle vertex to the hypotenuse. This altitude has length $h = \frac{a \cdot a}{a\sqrt{2}} = \frac{a^2}{a\sqrt{2}} = \frac{a}{\sqrt{2}} = \frac{a\sqrt{2}}{2}$. Setting this equal to $4\sqrt{2}$: $\frac{a\sqrt{2}}{2} = 4\sqrt{2}$, so $a\sqrt{2} = 8\sqrt{2}$, giving $a = 8$. The area is $\frac{1}{2} \times 8 \times 8 = \boxed{32}$.”




Figure 2: Early tokens encode strategy; later tokens merely execute. *Top:* example showing first 100 tokens (blue) establish the approach while last 100 tokens (orange) execute computation. *Bottom left:* prefix continuation—teacher performance degrades to student baseline by 300 tokens. *Bottom right:* best-checkpoint accuracy per 100-token region. First-100 yields +14.9pp; middle-100 and last-100 hurt.

approach. If the student’s strategy is suboptimal, the later-position signal teaches the student to better execute a suboptimal approach—rather than to choose a better one.

2.2 Per-Position KL Distribution

To investigate the distillation signal, we generate 59,936 on-policy trajectories (mean length 386 tokens) from the undistilled student and compute per-position KL divergence against the teacher. Positions are 0-indexed throughout (position 0 is the first generated token).

Figure 1 shows two complementary views of the signal distribution. The reverse KL divergence (left) reveals that the distillation signal is massively front-loaded: the first 50 tokens ($KL = 2.09$) carry $8.9\times$ more signal than the 900–1000 range ($KL = 0.24$), with a steep initial decay followed by a long, slowly declining tail beyond position 200. The teacher confidence analysis (right) explains why: at early positions, the teacher’s distribution has higher entropy (0.26 nats) and lower agreement with the student (79%). By position 1000, entropy drops to 0.12 and agreement rises to 93%—the teacher is no longer providing independent guidance, but merely confirming the student’s choices.

This is a direct consequence of the conditional scoring problem (§2.1): once the student’s prefix determines the reasoning approach, the teacher’s conditional distribution collapses to near-agreement. The residual disagreement at late positions is dominated by formatting conventions (LaTeX delimiters, code patterns) rather than reasoning content.

Strategy vs. execution. Token-type analysis (Appendix D) reveals that early high-KL tokens are predominantly *planning* tokens (“To”, “First”, “Let”)—reflecting strategy disagreement—while late high-KL tokens are formatting conventions. Figure 2 illustrates this with a concrete example: the first 100 tokens establish the solution approach, while the last 100 tokens execute the computation.

2.3 Test-Time Validation: Prefix Continuation

We directly test whether early tokens determine response quality via prefix continuation (Figure 2, left). With a 100-token student prefix, the teacher achieves 62.70% avg@4—only 2.6pp below its baseline. At 300 tokens, performance falls to 51.75%, nearly matching the student baseline: the student’s choices become entrenched. The rapid degradation quantifies the *window of effective influence*: the first ~ 100 tokens are where reasoning strategy is still malleable.

2.4 Training-Time Validation: Position Region Ablation

The analysis above predicts that early-position training should be most effective and later-position training should be harmful. We test this by training on different 100-token regions (Figure 2, right).

First-100 improves avg@4 by +14.9pp and outperforms full-sequence training. Middle-100 and last-100 both perform *below the undistilled baseline*—later-position signals are not merely uninformative, they are actively harmful. Random-100 (63.05%) outperforms full-seq (62.35%), further confirming that including later-position tokens is a net negative.

2.5 Training Instability as a Consequence

Figure 3 shows that full-sequence distillation degrades with continued training on coding (HumanEval drops from 40.2% to 26.8%, -33%), while positional distillation remains stable. This is a direct consequence of the signal quality analysis: over many training steps, noisy gradients from later positions accumulate and distort the student’s parameters.

2.6 Summary of Findings

Our analysis reveals a coherent picture: (1) the teacher’s signal is strong and independent at early positions, where reasoning strategy is being determined; (2) the signal decays to near-agreement at later positions, where the teacher is conditioned on the student’s prefix; (3) training on later positions is harmful, and full-sequence degradation is a consequence of this harm. This motivates restricting distillation to early positions.

3 Method: Positional Distillation

3.1 Formulation

In standard on-policy reverse KL distillation, the student p_θ generates a response $\mathbf{x} = (x_1, \dots, x_T)$ conditioned on a prompt, and the loss is computed against the teacher q_ϕ :

$$\mathcal{L}_{\text{full}} = \mathbb{E}_{\mathbf{x} \sim p_\theta} \left[\sum_{t=1}^T \text{KL}(q_\phi(\cdot | \mathbf{x}_{<t}) \parallel p_\theta(\cdot | \mathbf{x}_{<t})) \right]. \quad (1)$$

Positional distillation restricts the loss to the first N response tokens:

$$\mathcal{L}_{\text{pos}} = \mathbb{E}_{\mathbf{x} \sim p_\theta} \left[\sum_{t=1}^N \text{KL}(q_\phi(\cdot | \mathbf{x}_{<t}) \parallel p_\theta(\cdot | \mathbf{x}_{<t})) \right], \quad (2)$$

where N is a fixed position cutoff ($N \ll T$ in practice). If the student generates an EOS token before position N , loss is computed only up to EOS.

In practice, we both restrict the loss *and* truncate the rollout to N tokens. This couples two benefits: loss masking avoids low-quality tail positions (stability), while rollout truncation reduces generation cost (efficiency). Since autoregressive generation is the dominant cost in on-policy KD, truncating from $T \sim 400$ to $N = 100$ tokens yields $20\text{--}35\times$ wall-clock speedup (Table 4).

Table 1: Main results. Best performance across training steps for each configuration. **Bold:** best per teacher. Full-seq results use the best step before degradation. Time: seconds per training step (gen + score + train). Gemma uses gemma-2-2b-it student $\rightarrow$ gemma-3-4b-it teacher.

Teacher	Method	MATH-500 avg@4	HumanEval pass@1	MBPP pass@1	BFCL full_acc	Time (s/step)
<i>Student: Qwen2.5-Math-1.5B (LoRA)</i>						
—	Baseline	50.95	32.93	52.12	2.70	—
Qwen3-1.7B	Full-seq	62.35	40.24	52.65	—	~ 280
	Ours ($N=100$)	65.85	42.07	49.21	61.30	~ 8
Qwen3-4B	Full-seq	67.45	39.63	50.79	34.83	~ 210
	Ours ($N=100$)	68.95	43.29	52.38	45.80	~ 8
<i>Student: gemma-2-2b-it (LoRA) — cross-family</i>						
—	Baseline	13.45	23.78	40.48	73.50	—
gemma-3-4b-it	Full-seq	11.70 [§]	—	—	—	~ 210
	Ours ($N=100$)	27.20	31.10	50.53	82.50	~ 8

[§] Gemma full-seq degrades below baseline; other tasks not evaluated.

4 Generalization Experiments

4.1 Setup

Models. We evaluate across multiple student-teacher pairs. Our primary student is Qwen2.5-Math-1.5B [Qwen Team, 2024], paired with teachers of increasing scale: Qwen3-1.7B, Qwen3-4B, and Qwen3-8B [Qwen Team, 2025]. To test cross-family generality, we also evaluate with gemma-2-2b-it (student) and gemma-3-4b-it (teacher), which use different tokenizers requiring cross-tokenizer re-tokenization and shared-vocabulary KL computation.

Training. All experiments use LoRA [Hu et al., 2022] ($r=32$, $\alpha=64$, all linear layers, $\text{lr} = 5 \times 10^{-5}$) with reverse KL divergence loss. Temperature 0.7 for on-policy generation. Each training step processes a batch of 16 problems with 1 rollout per problem ($n_{\text{samples}}=1$, batch size 16). We train for 200 steps on math and function calling, 400 steps on coding, saving checkpoints every 50 steps. Training data: 3,200 math problems, 6,400 coding problems (CodeUltraFeedback), and 3,200 function calling examples (glaive-function-calling-v2). Our method uses $N=100$ unless otherwise specified.

Evaluation. MATH-500 [Hendrycks et al., 2021, Lightman et al., 2023] with $n=4$ samples at temperature 0.7 (reporting avg@4). HumanEval/HumanEval+ and MBPP/MBPP+ [Chen et al., 2021, Austin et al., 2021, Liu et al., 2024] at temperature 0.0 (reporting pass@1). BFCL [Yan et al., 2024] with 600 problems (reporting full accuracy: correct function name and arguments).


4.2 Main Results

Table 1 presents the unified main results across tasks, teacher scales, and model families. We report the best performance achieved across all training steps for each configuration.

Results. With the Qwen3-1.7B teacher, positional distillation ($N=100$) matches or exceeds full-sequence distillation on 3 of 4 benchmarks while achieving $35\times$ speedup. Scaling to the 4B teacher improves quality further: 68.95% avg@4 on math (+3.1pp over 1.7B), with positional distillation outperforming full-seq on all metrics. The Gemma cross-family pair demonstrates that positional distillation transfers across model families with incompatible tokenizers, doubling math accuracy (13.45% $\rightarrow$ 27.20%) and substantially improving all other metrics. Additional teacher scales (8B) are in Appendix B.

Table 2: Training efficiency. Per-step time and peak GPU memory on a single A6000 (48GB).

Teacher	Method	Time/step	Peak GPU mem	Total (200 steps)
Q3-1.7B	Full-seq	~280s	43 GB (vLLM)	~16 hours
	Ours ($N=100$)	~8s	8 GB (HF)	~27 min
Q3-4B	Full-seq	~210s	46 GB (vLLM)	~12 hours
	Ours ($N=100$)	~8s	12 GB (HF)	~27 min

**Figure 3: Full-sequence distillation degrades on coding.** On HumanEval, full-seq drops from 40.2% to 26.8% (−33%) over 400 training steps, while first-100 positional distillation remains stable above baseline.

Cascade effect. A natural concern is that distilling only the first N tokens might only improve those positions. We find the opposite: early-token distillation creates a *cascade effect* that reshapes the entire response. Because the early tokens determine the reasoning strategy, improving them causes the subsequent execution tokens to change accordingly—even though those later positions receive zero distillation signal. Concretely, late tokens change just as much as early tokens (Jaccard similarity to baseline = 0.251 for both halves), and pos-200 reduces KL by 50–60% within the trained range (positions 0–200) and by 30–40% in the *untrained* range (positions 200–400). This is consistent with the strategy-then-execution structure identified in Section 2.2: better strategy decisions at early positions naturally lead to better execution at later positions.

4.3 Training Stability

As shown in Figure 3 and discussed in Section 2.5, full-sequence distillation degrades over training while positional distillation remains stable. This is one of the three key advantages of positional distillation: better performance, no degradation, and dramatically faster training.

4.4 Efficiency

Table 4 shows that positional distillation achieves 26–35 $\times$ wall-clock speedup. The dominant cost in full-seq is autoregressive generation (~ 280 s/step for sequences averaging ~ 1000 tokens); positional distillation generates only 100 tokens (~ 8 s/step). Positional distillation also requires significantly less GPU memory because it does not need a vLLM/SGLang inference server—100 tokens can be generated via standard HF generate, enabling training on a single consumer GPU.




Figure 4: Position limit sweep on MATH-500. Best avg@4 across training steps for each N . Performance increases monotonically and surpasses full-seq (dashed orange) at $N \geq 150$. Even $N=5$ provides +5.5pp over baseline (dashed gray).

5 Ablation Experiments

5.1 Position Limit Sweep

How sensitive is positional distillation to the choice of N ? Figure 4 shows best avg@4 across $N \in \{5, 10, 20, 50, 100, 150, 200\}$ on MATH-500.

The results reveal a robust sweet spot at $N = 50\text{--}200$, where performance saturates. All positional variants outperform full-sequence distillation (65.55%), with $N=200$ achieving the best result (66.75%). Even $N=5$ provides +5.5pp over baseline, confirming that the very first tokens carry substantial distillation signal.

5.2 Token Selection Strategy Comparison

A natural alternative to positional selection is *information-theoretic* selection: choose the k tokens with the highest KL divergence or entropy, regardless of position. We compare four selection strategies, each selecting exactly 100 tokens from full-length student generations (Figure 5).

The results reveal a striking pattern: **the strategies that capture the most signal perform worst.** Top-KL-100 captures 95.6% of total KL divergence yet achieves only 58.60%—the worst of all methods. The key difference is *position*: high-KL and high-entropy tokens are concentrated at later positions (mean position 240–313), while prefix tokens span positions 0–99. This confirms that **position is the most important criterion for token selection in on-policy distillation.**

5.3 LoRA vs. Full Fine-Tuning

We compare LoRA ($r=32$, $\text{lr} = 5 \times 10^{-5}$) against full fine-tuning ($\text{lr} = 5 \times 10^{-6}$) on both math and coding tasks with $N=100$.

LoRA substantially outperforms full fine-tuning on both tasks (+9.7pp on math, +5.5pp on HumanEval for pos-100). This is consistent with the low-rank regularization preventing the student from overfitting to the distillation signal—particularly important given the small model size (1.5B parameters) and limited training data (3,200 problems). Full fine-tuning’s advantage on MBPP (54.2% vs. 49.2%) suggests the methods may be complementary across benchmarks.

Additional student-teacher pairs. We evaluate with Qwen3-1.7B and Qwen3-4B as students (with 4B and 8B teachers). Distillation is most effective when the student has




Figure 5: Token selection strategies ($k=100$, best avg@4). Prefix (first 100 tokens) outperforms all alternatives. Top-KL and top-entropy strategies select tokens with the highest signal concentration (95%+ of total KL/entropy) but these tokens are at later positions (mean pos. 240–313) where the teacher signal is harmful. Random-100 outperforms all “smart” strategies except prefix by statistically avoiding later positions.

Table 3: LoRA vs. full fine-tuning. Best performance across training steps. Qwen2.5-Math-1.5B $\rightarrow$ Qwen3-1.7B, $n=1$, $bs=16$.

Task	Method	LoRA	Full FT
MATH-500 avg@4	Pos-100	65.85%	56.20%
	Full-seq	62.35%	58.20%
HumanEval pass@1	Pos-100	42.1%	36.6%
	Full-seq	40.2%	36.6%

significant room for improvement: Qwen3-1.7B gains +13.9pp on BFCL from a 4B teacher. Stronger students show diminishing returns (Appendix B).

6 Related Work

Knowledge Distillation for Language Models. Knowledge distillation [Hinton et al., 2015] transfers knowledge from a teacher model to a smaller student by training on soft targets. Kim and Rush [2016] extended this to sequence models with word-level and sequence-level distillation objectives. For large language models, the choice of divergence measure and data distribution has proven critical. Gu et al. [2024] demonstrated that reverse KL divergence is preferable for generative LLM distillation, as it avoids the student overestimating low-probability teacher modes. Agarwal et al. [2024] introduced Generalized Knowledge Distillation (GKD), showing that on-policy generation—where the student generates sequences and the teacher provides soft labels—substantially outperforms supervised (off-policy) distillation. Ko et al. [2024] proposed skew KL divergence and adaptive off-policy approaches for more efficient distillation, while Wen et al. [2023] generalized sequence-level KD to arbitrary f-divergences. Wu et al. [2025] challenged the conventional mode-seeking vs. mean-seeking interpretation of reverse and forward KL in the LLM setting, finding that both divergences converge given sufficient training but focus on different distribution regions in early epochs. More recently, Xu et al. [2025] introduced Speculative Knowledge Distillation, which uses interleaved teacher-student sampling inspired by speculative decoding, filtering tokens the teacher is unlikely to produce. Our work builds on the on-policy reverse KL paradigm of Agarwal et al. [2024] and Gu et al. [2024], but introduces position-based loss

computation that fundamentally changes which tokens receive gradient signal.

Token-Level Importance in Distillation and Reasoning. A growing body of work demonstrates that not all tokens contribute equally to learning. Wang et al. [2025] showed that only $\sim 20\%$ of chain-of-thought tokens have high entropy, and these “forking tokens” steer reasoning paths—applying RL gradients to only these tokens matches full-gradient performance. In the RL fine-tuning setting, Vassoyan et al. [2025] demonstrated that uniform KL penalties across all token positions prevent exploration on “critical tokens” where the pretrained model is most uncertain, and proposed entropy-weighted KL relaxation that allows the model to deviate from its prior precisely where it needs to learn. Li et al. [2025] discovered a “preplan-and-anchor” mechanism in LLM attention patterns, where models first perform long-range contextual reference before producing semantic anchor tokens, enabling targeted credit assignment. Singh and Hakkani-Tür [2026] proposed greedy pruning of reasoning chains and showed that LLMs internally encode token-level functional importance for answer generation.

In the distillation setting specifically, several concurrent works have explored token selection. Huang et al. [2025] proposed SelecTKD, a propose-and-verify framework where the teacher verifies student-proposed tokens, applying full loss to accepted tokens and masking rejected ones. Xie et al. [2026] introduced AdaKD, which adaptively adjusts token-level temperature based on learning stability. Tavor et al. [2026] systematically disentangled selective KD along position, class, and sample axes, proposing student-entropy-guided position selection (SE-KD) that reduces wall time by 70% without sacrificing accuracy. Kim and Baek [2026] combined entropy-based token selection with preference ranking in a dual distillation framework.

These works collectively establish that not all tokens deserve equal treatment during training. Our approach differs in using *absolute position* as the sole selection criterion, requiring no per-token computation of entropy, difficulty, or verification scores. This simplicity is motivated by our empirical finding that KL divergence between teacher and student concentrates at early positions corresponding to reasoning strategy selection. Furthermore, unlike prior token-selection methods that aim primarily at efficiency or robustness, we show that positional loss also eliminates training instability—a qualitative benefit that token-level adaptive methods do not address.

Training Stability and Sequence-Level Issues. Full-sequence on-policy distillation is known to suffer from instability. Holtzman et al. [2020] identified that likelihood-based training can produce bland, repetitive text—a form of neural text degeneration. Shumailov et al. [2024] demonstrated that iteratively training generative models on their own outputs leads to model collapse, where distribution tails disappear. In on-policy KD specifically, Jang et al. [2026] proposed Veto, a geometric bridge in logit space that suppresses harmful gradients on low-confidence tokens, directly addressing the instability we also observe. Our approach is complementary: rather than modifying the target distribution, we restrict the loss scope to early positions where teacher-student agreement is naturally highest, achieving stability through a simpler mechanism.

Several works have studied how sequence length affects distillation quality. Chen et al. [2025] showed that training on only the first 50% of tokens in supervised reasoning distillation retains $\sim 91\%$ of full-sequence performance, establishing a computation-quality tradeoff. Yan et al. [2026] addressed exposure bias and capacity misalignment in long chain-of-thought distillation through temperature-scheduled learning. Our work extends the truncation insight to the on-policy setting: we demonstrate that positional loss masking on student-generated sequences achieves equal or better performance than full-sequence loss, while providing the additional benefit of eliminating degenerate repetition patterns that are unique to on-policy training.

Curriculum Learning and Sequence Scheduling. Curriculum learning [Bengio et al., 2009] improves training by gradually increasing task difficulty. Bengio et al. [2015] applied this

principle to sequence models through scheduled sampling, transitioning from teacher-forced to free-running generation during training. In the LLM pretraining setting, Li et al. [2022] showed that long sequences at training start cause extreme gradient variance, and sequence length warmup stabilizes training while enabling larger batch sizes. Pouransari et al. [2024] extended this with variable sequence length curriculum for LLM pretraining, achieving up to $3\times$ faster convergence. Our progressive position schedule—linearly increasing the position limit from 1 to 200 over training—applies the curriculum principle to the distillation loss scope rather than to the input sequence length, and represents a novel form of curriculum specific to knowledge distillation.

Reasoning Distillation. The distillation-vs-RL paradigm for reasoning has been established by DeepSeek-R1 [DeepSeek-AI, 2025], which showed that distilled models can outperform RL-trained ones in certain settings. Chain-of-thought prompting [Wei et al., 2022] enables LLMs to solve complex problems via intermediate reasoning, and distilling CoT capabilities has become a key application of KD. Our work contributes to this landscape by demonstrating that the effectiveness of distillation can be significantly improved through principled control of the loss computation scope.

7 Discussion

The conditional scoring problem is fundamental. Our analysis reveals that a core issue in on-policy distillation is often overlooked: the teacher’s supervision at position t is not an unconditional assessment of token quality, but a *conditional* assessment given the student’s prefix $y_{<t}$. At early positions, this conditioning is minimal (the prefix is mostly the prompt), so the teacher provides genuinely independent guidance on reasoning strategy. At later positions, the teacher is locked into the student’s chosen approach and can only score execution details within that approach. This “conditional scoring problem” explains both why early-position signal is valuable (it transfers strategy) and why later-position signal is harmful (it reinforces potentially suboptimal approaches).

Broader implications. The signal quality degradation we identify may extend beyond distillation to other on-policy optimization methods that use token-level rewards, including RLHF and DPO. Whenever a reward model or reference model scores tokens conditioned on a learner-generated prefix, the same conditional dependency applies: later-position rewards are increasingly shaped by the learner’s choices rather than representing independent quality assessments. Investigating positional signal quality in these settings is a promising direction.

Limitations. Our signal quality analysis focuses on the Qwen model family; extending to other families and larger scales would validate generality. We evaluate only reverse KL divergence; forward KL or other f -divergences may show different positional patterns. The optimal N varies by task (50–200 works broadly); adaptive per-example N selection could be more robust.

Conclusion. We show that in on-policy knowledge distillation, the teacher’s per-token supervision quality is fundamentally non-uniform: early tokens carry strong, independent signal about reasoning strategy, while later tokens provide weak, dependent signal that is actively harmful when optimized. This insight motivates Positional Distillation—restricting KD loss to the first N response tokens—which matches or exceeds full-sequence distillation across math, coding, and function calling tasks while eliminating training instability and achieving $20\text{--}35\times$ speedup. Teaching the student how to begin is sufficient to improve the entire response.

References

- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated mistakes. In *International Conference on Learning Representations*, 2024.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Samy Bengio, Oriol Vinyals, Navdeep Jaitly, and Noam Shazeer. Scheduled sampling for sequence prediction with recurrent neural networks. In *Advances in Neural Information Processing Systems*, 2015.
- Yoshua Bengio, Jérôme Louradour, Ronan Collobert, and Jason Weston. Curriculum learning. In *Proceedings of the 26th Annual International Conference on Machine Learning*, pages 41–48, 2009.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Wei-Rui Chen, Vignesh Kothapalli, Ata Fatahibaarzi, Hejian Sang, Shao Tang, Qingquan Song, Zhipeng Wang, and Muhammad Abdul-Mageed. Distilling the essence: Efficient reasoning distillation via sequence truncation. *arXiv preprint arXiv:2512.21002*, 2025.
- DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge distillation of large language models. In *International Conference on Learning Representations*, 2024.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In *Advances in Neural Information Processing Systems*, 2021.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. In *NIPS Deep Learning and Representation Learning Workshop*, 2015.
- Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. In *International Conference on Learning Representations*, 2020.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- Haiduo Huang, Jiangcheng Song, Yadong Zhang, and Pengju Ren. SelecTKD: Selective token-weighted knowledge distillation for LLMs. *arXiv preprint arXiv:2510.24021*, 2025.
- Ijun Jang, Jewon Yeom, Juan Yeo, Hyunggu Lim, and Taesup Kim. Stable on-policy distillation through adaptive target reformulation. *arXiv preprint arXiv:2601.07155*, 2026.
- Minsang Kim and Seung Jun Baek. Explain in your own words: Improving reasoning via token-selective dual knowledge distillation. *arXiv preprint arXiv:2603.13260*, 2026.
- Yoon Kim and Alexander M. Rush. Sequence-level knowledge distillation. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 1317–1327, 2016.

- Jongwoo Ko, Sungnyun Kim, Tianyi Chen, and Se-Young Yun. DistiLLM: Towards streamlined distillation for large language models. In *International Conference on Machine Learning*, 2024.
- Conglong Li, Minjia Zhang, and Yuxiong He. The stability-efficiency dilemma: Investigating sequence length warmup for training GPT models. In *Advances in Neural Information Processing Systems*, 2022.
- Yang Li, Zhichen Dong, Yuhan Sun, Weixun Wang, Shaopan Xiong, Yijia Luo, Jiashun Liu, Han Lu, Jiamang Wang, Wenbo Su, Bo Zheng, and Junchi Yan. Attention illuminates LLM reasoning: The preplan-and-anchor rhythm enables fine-grained policy optimization. *arXiv preprint arXiv:2510.13554*, 2025.
- Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. *arXiv preprint arXiv:2305.20050*, 2023.
- Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. Is your code generated by ChatGPT really correct? rigorous evaluation of large language models for code generation. *Advances in Neural Information Processing Systems*, 2024.
- Kevin Lu, Tianjun Chen, Caglar Gulcehre, et al. On-policy distillation. *Thinking Machines Lab Blog*, 2025. <https://thinkingmachines.ai/blog/on-policy-distillation/>.
- Hadi Pouransari, Chun-Liang Li, Jen-Hao Rick Chang, Pavan Kumar Anasosalu Vasu, Cem Koc, Vaishaal Shankar, and Oncel Tuzel. Dataset decomposition: Faster LLM training with variable sequence length curriculum. In *Advances in Neural Information Processing Systems*, 2024.
- Qwen Team. Qwen2.5-math technical report: Toward mathematical expert model via self-improvement. *arXiv preprint arXiv:2409.12122*, 2024.
- Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- Ilya Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. AI models collapse when trained on recursively generated data. *Nature*, 631(8022): 755–759, 2024.
- Janvijay Singh and Dilek Hakkani-Tür. Do LLMs encode functional importance of reasoning tokens? *arXiv preprint arXiv:2601.03066*, 2026.
- Almog Tavor, Itay Ebenspanger, Neil Cnaan, and Mor Geva. Rethinking selective knowledge distillation. *arXiv preprint arXiv:2602.01395*, 2026.
- Jean Vassoyan, Nathanael Beau, and Roman Plaud. Ignore the KL penalty! boosting exploration on critical tokens to enhance RL fine-tuning. In *Findings of the Association for Computational Linguistics: NAACL 2025*, 2025.
- Shenzhi Wang, Le Yu, Chang Gao, Chujie Zheng, Shixuan Liu, Rui Lu, Kai Dang, Xionghui Chen, Jianxin Yang, Zhenru Zhang, Yuqiong Liu, An Yang, Andrew Zhao, Yang Yue, Shiji Song, Bowen Yu, Gao Huang, and Junyang Lin. Beyond the 80/20 rule: High-entropy minority tokens drive effective reinforcement learning for LLM reasoning. In *Advances in Neural Information Processing Systems*, 2025.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*, 2022.

Yuqiao Wen, Zichao Li, Wenyu Du, and Lili Mou. f-divergence minimization for sequence-level knowledge distillation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10817–10834, 2023.

Taiqiang Wu, Chaofan Tao, Jiahao Wang, Runming Yang, Zhe Zhao, and Ngai Wong. Rethinking Kullback-Leibler divergence in knowledge distillation for large language models. In *Proceedings of the International Conference on Computational Linguistics (COLING)*, 2025.

Xurong Xie, Zhucun Xue, Jiafu Wu, Jian Li, Yabiao Wang, Xiaobin Hu, Yong Liu, and Jiangning Zhang. LLM-oriented token-adaptive knowledge distillation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2026.

Wenda Xu, Rujun Han, Zifeng Wang, Long T. Le, Dhruv Madeka, Lei Li, William Yang Wang, Rishabh Agarwal, Chen-Yu Lee, and Tomas Pfister. Speculative knowledge distillation: Bridging the teacher-student gap through interleaved sampling. In *International Conference on Learning Representations*, 2025.

Fanjia Yan, Huanzhi Mao, Charlie Ji, Tianjun Zhang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. Berkeley function calling leaderboard. *arXiv preprint arXiv:2402.15671*, 2024.

Shaotian Yan, Kaiyuan Liu, Chen Shen, Bing Wang, Sinan Fan, Jun Zhang, Yue Wu, Zheng Wang, and Jieping Ye. Distribution-aligned sequence distillation for superior long-CoT reasoning. *arXiv preprint arXiv:2601.09088*, 2026.

Wenkai Yang, Weijie Liu, Ruobing Xie, Kai Yang, Saiyong Yang, and Yankai Lin. Learning beyond teacher: Generalized on-policy distillation with reward extrapolation. *arXiv preprint arXiv:2602.12125*, 2026.

A Training Efficiency

Table 4: Training efficiency. Per-step wall-clock time on a single A6000 (48GB). Student: Qwen2.5-Math-1.5B, LoRA, bs=16.

Teacher	Method	Gen (s)	Score (s)	Train (s)	Total (s)
Q3-1.7B	Full-seq	100–731	3–10	5–12	~280
	Ours ($N=100$)	~5	~1	~2	~8
Q3-4B	Full-seq	97–733	3–9	6–10	~210
	Ours ($N=100$)	~5	~1	~2	~8
Q3-8B	Full-seq	96–230	7–11	1–6	~170 [†]
	Ours ($N=100$)	~5	~1	~2	~8

[†] Frequent OOMs; 48GB insufficient for 8B teacher + vLLM + student.

B Additional Student-Teacher Pairs

C Full Experimental Results

C.1 Math Results: Per-Step Performance

Table 6 presents per-step results for the primary math experiments (LoRA, $n=1, 3, 200$ problems).

Table 5: Additional student-teacher pairs. All use pos-100, LoRA, 3,200 problems.

Student	Teacher	MATH-500 avg@4	BFCL full_acc	HumanEval pass@1
Qwen3-1.7B	— (baseline)	69.20%	59.80%	39.63%
Qwen3-1.7B	Qwen3-4B	69.20%	73.70%	44.51%
Qwen3-1.7B	Qwen3-8B	69.15%	74.00%	40.24%
Qwen3-4B	— (baseline)	77.95%	73.00%	73.17%
Qwen3-4B	Qwen3-8B	77.05%	76.80%	71.34%

Table 6: MATH-500 per-step results. LoRA, $n=1$, $bs=16$, 3,200 problems. Baseline: 50.95% avg@4.

Method	Metric	Step 50	Step 100	Step 150	Step 200
Pos-50	avg@4	62.35	66.05	66.65	64.85
	maj@4	69.40	72.00	71.00	71.20
	pass@4	77.20	79.40	81.00	79.60
Pos-100	avg@4	63.75	64.45	65.15	65.85
	maj@4	70.00	68.40	69.60	70.80
	pass@4	79.80	78.40	80.20	79.80
Pos-150	avg@4	65.35	66.65	65.30	65.75
	maj@4	66.80	67.00	66.30	67.30
	pass@4	79.00	81.00	78.20	80.00
Pos-200	avg@4	66.05	64.65	65.10	65.55
	maj@4	71.20	68.40	70.00	71.20
	pass@4	81.00	79.80	80.60	80.60

Table 7: Complete MATH-500 results, LoRA, $n=16$, $bs=16$. Best per configuration in **bold**.

Config	Step 50	Step 100	Step 150	Step 200
<i>avg@4</i>				
Pos-5	55.15	55.65	56.50	56.30
Pos-10	59.50	56.40	56.30	56.00
Pos-20	60.20	57.35	56.40	56.55
Pos-50	62.45	62.15	60.50	61.90
Pos-100	61.15	63.55	64.25	63.85
Pos-150	63.55	65.70	64.80	64.40
Pos-200	63.25	66.75	64.55	65.25
Progressive	61.10	62.05	61.95	62.30
Full-seq [†]	65.60	46.30	49.70	43.80
Full-seq (first-boxed)	65.55	64.95	65.55	64.95
<i>pass@4</i>				
Pos-5	74.0	74.8	76.0	73.4
Pos-10	75.8	73.2	75.4	74.2
Pos-20	77.0	76.0	73.4	75.8
Pos-50	77.0	78.4	77.6	78.6
Pos-100	77.6	79.2	80.2	80.2
Pos-150	79.6	77.4	80.0	79.4
Pos-200	79.6	80.2	79.4	79.2
Progressive	77.2	77.8	78.2	78.8
Full-seq	80.0	72.2	77.0	74.2

[†] Last-boxed extraction; degrades due to repetition.

Table 8: Complete coding results, LoRA. HumanEval (HE) and MBPP pass@1.

Config	Metric	s50	s100	s150	s200	s250	s300	s350	s400
Pos-50	HE	37.8	39.0	39.6	41.5	40.2	40.9	42.1	40.9
	MBPP	52.1	48.9	46.0	46.0	47.6	46.3	46.6	46.3
Pos-100	HE	37.2	39.0	42.1	37.8	39.0	37.8	37.8	38.4
	MBPP	52.1	49.7	49.2	49.2	48.4	49.2	47.9	48.9
Pos-150	HE	36.6	35.4	36.6	39.0	41.5	39.6	38.4	37.2
	MBPP	51.3	50.0	48.4	50.5	49.5	50.3	49.7	48.7
Full-seq	HE	40.2	31.7	32.3	32.9	27.4	28.0	26.8	26.8
	MBPP	52.6	48.9	49.5	48.1	47.9	47.1	48.4	47.6

Table 9: Function calling results (BFCL), LoRA. Name accuracy / Full accuracy / Parse rate. Best full_acc in **bold**.

Method	Best Step	Name Acc	Full Acc	Parse Rate
Baseline	—	9.70%	2.70%	24.20%
Teacher (Qwen3-1.7B)	—	75.30%	54.00%	75.30%
Pos-50	200	95.20%	57.20%	98.30%
Pos-100	100	86.20%	61.30%	91.30%
Pos-150	200	88.70%	61.50%	92.50%
Pos-200	200	80.80%	54.50%	90.20%
Full-seq	100	81.00%	58.20%	86.70%

Table 10: Scaling results. All use pos-100, LoRA, 3,200 problems. Best across training steps. Baseline performance shown for reference.

Student → Teacher	Math avg@4	BFCL full_acc	HE pass@1	MBPP+ pass@1
M-1.5B baseline	50.95%	2.70%	—	—
M-1.5B → Q3-1.7B	65.85%	61.30%	42.07%	44.44%
M-1.5B → Q3-4B	68.95%	45.80%	43.29%	47.62%
M-1.5B → Q3-8B	67.85%	57.00%	41.46%	44.44%
Q3-1.7B baseline	69.20%	59.80%	39.63%	44.71%
Q3-1.7B → Q3-4B	69.20%	73.70%	44.51%	47.62%
Q3-1.7B → Q3-8B	69.15%	74.00%	40.24%	48.15%
Q3-4B baseline	77.95%	73.00%	73.17%	55.56%
Q3-4B → Q3-8B	77.05%	76.80%	71.34%	56.35%

C.2 Math Results: Earlier Series ($n=16$, 200 problems)

C.3 Coding Results

C.4 Function Calling Results

C.5 Scaling Experiment Results

C.6 Full-Sequence Degradation Details

Table 11: Boxed repetition progression in full-sequence LoRA distillation on MATH-500 ($n=16$, $bs=16$).

Statistic	Step 50	Step 100	Step 150	Step 200
Multi-boxed responses	2.9%	91.5%	80.5%	89.1%
Avg <code>\boxed{}</code> per response	1.0	88.1	54.1	58.0
Repetitive responses	5.9%	78.8%	58.3%	77.5%
Avg response length (chars)	1,388	4,023	3,585	4,558
Official avg@4	65.60%	46.30%	49.70%	43.80%
Relaxed avg@4 (GT in text)	74.60%	70.90%	71.90%	70.30%

D Token Classification Methodology

We classify each token into six categories based on string matching:

1. **planning**: Reasoning keywords (“To”, “Let”, “First”, “Step”, “We”, “Given”, “Therefore”, “Thus”, “Since”).
2. **structural**: Punctuation, whitespace, formatting tokens.
3. **math_number**: Digits (0–9).
4. **math_operator**: Arithmetic operators (+, −, ×, /, =).
5. **math_latex**: LaTeX delimiters ($\backslash$ (, $\backslash$ [).
6. **continuation**: All others.

Table 12: Mean KL by token category and position range.

Category	0–4	5–19	20–49	50–99	100–199	200–499
planning	4.50	0.79	1.49	1.66	1.49	2.37
structural	3.26	1.46	1.60	0.93	0.60	0.86
math_number	1.49	0.60	0.74	0.28	0.17	0.13
math_operator	7.30	1.84	0.81	0.37	0.21	0.14
math_latex	8.84	9.23	6.50	4.95	2.97	1.87
continuation	1.94	1.12	1.19	0.89	0.70	0.48

Table 13: Top 20 highest-KL tokens (minimum 50 occurrences across 10,000 trajectories).

Rank	Token	Category	Count	Mean KL
1	“Solution”	planning	152	21.93
2	“Analysis”	continuation	125	16.49
3	\[	math_latex	7,152	13.21
4	“examines”	continuation	74	11.51
5	“He”	continuation	150	10.80
6	\(	math_latex	21,243	10.30
7	“First”	planning	1,706	9.98
8	“tests”	continuation	52	8.71
9	\	math_latex	82	8.68
10	“There”	continuation	201	8.28
11	“Therefore”	planning	4,913	7.95
12	“Identify”	continuation	1,345	7.78
13	“To”	planning	8,806	6.90
14	“When”	continuation	89	6.62
15	“This”	continuation	621	6.34
16	“Thus”	planning	2,547	6.24
17	“First” (space)	planning	259	6.10
18	“To” (space)	planning	1,174	5.37
19	“The”	planning	2,626	5.25
20	“Next”	planning	1,753	5.03